

# Operating System Labs

Yuanbin Wu  
CS@ECNU

# Operating System Labs

- Project 2 Due
  - 21:00, Nov. 12
- Project 2
  - group of 3
  - You now have 3 “late days”, but start early!
  - Oral test at week 12 (Nov. 27)
- We encourage using **git** for your group assignments (geeks' choice).

# Operating System Labs

- C Memory API
- Free Memory Management

# C Memory API

- Type of memory
  - Stack
  - Heap

# C Memory API

- Stack
  - Allocated / Deallocate automatically
  - By the compiler
  - Automatic memory

# C Memory API

- Stack

- Example (local variable)

```
void func()
{
    int x = 0;
    ...
}
```

- You only declare the variable
- Compiler will allocate it when someone calls the function
- Also will deallocate it when func returns

# C Memory API

- Heap
  - Allocated / Deallocate explicitly
  - By you, the programmer

# C Memory API

- Heap

- Example (malloc)

```
void func()
{
    int *ptr = (int*)malloc(sizeof(int));
    ...
}
```

- Both stack and heap allocation
  - When func returns,
    - Stack memory will be deallocated
    - Heap memory is still there



# C Memory API

- Stack and Heap

- Heap

- From low addr to high addr

- Stack

- From high addr to low addr


- Let's see

- mmap

- front-end for smap file proc file system

FFFFFFFF

Stack



Free

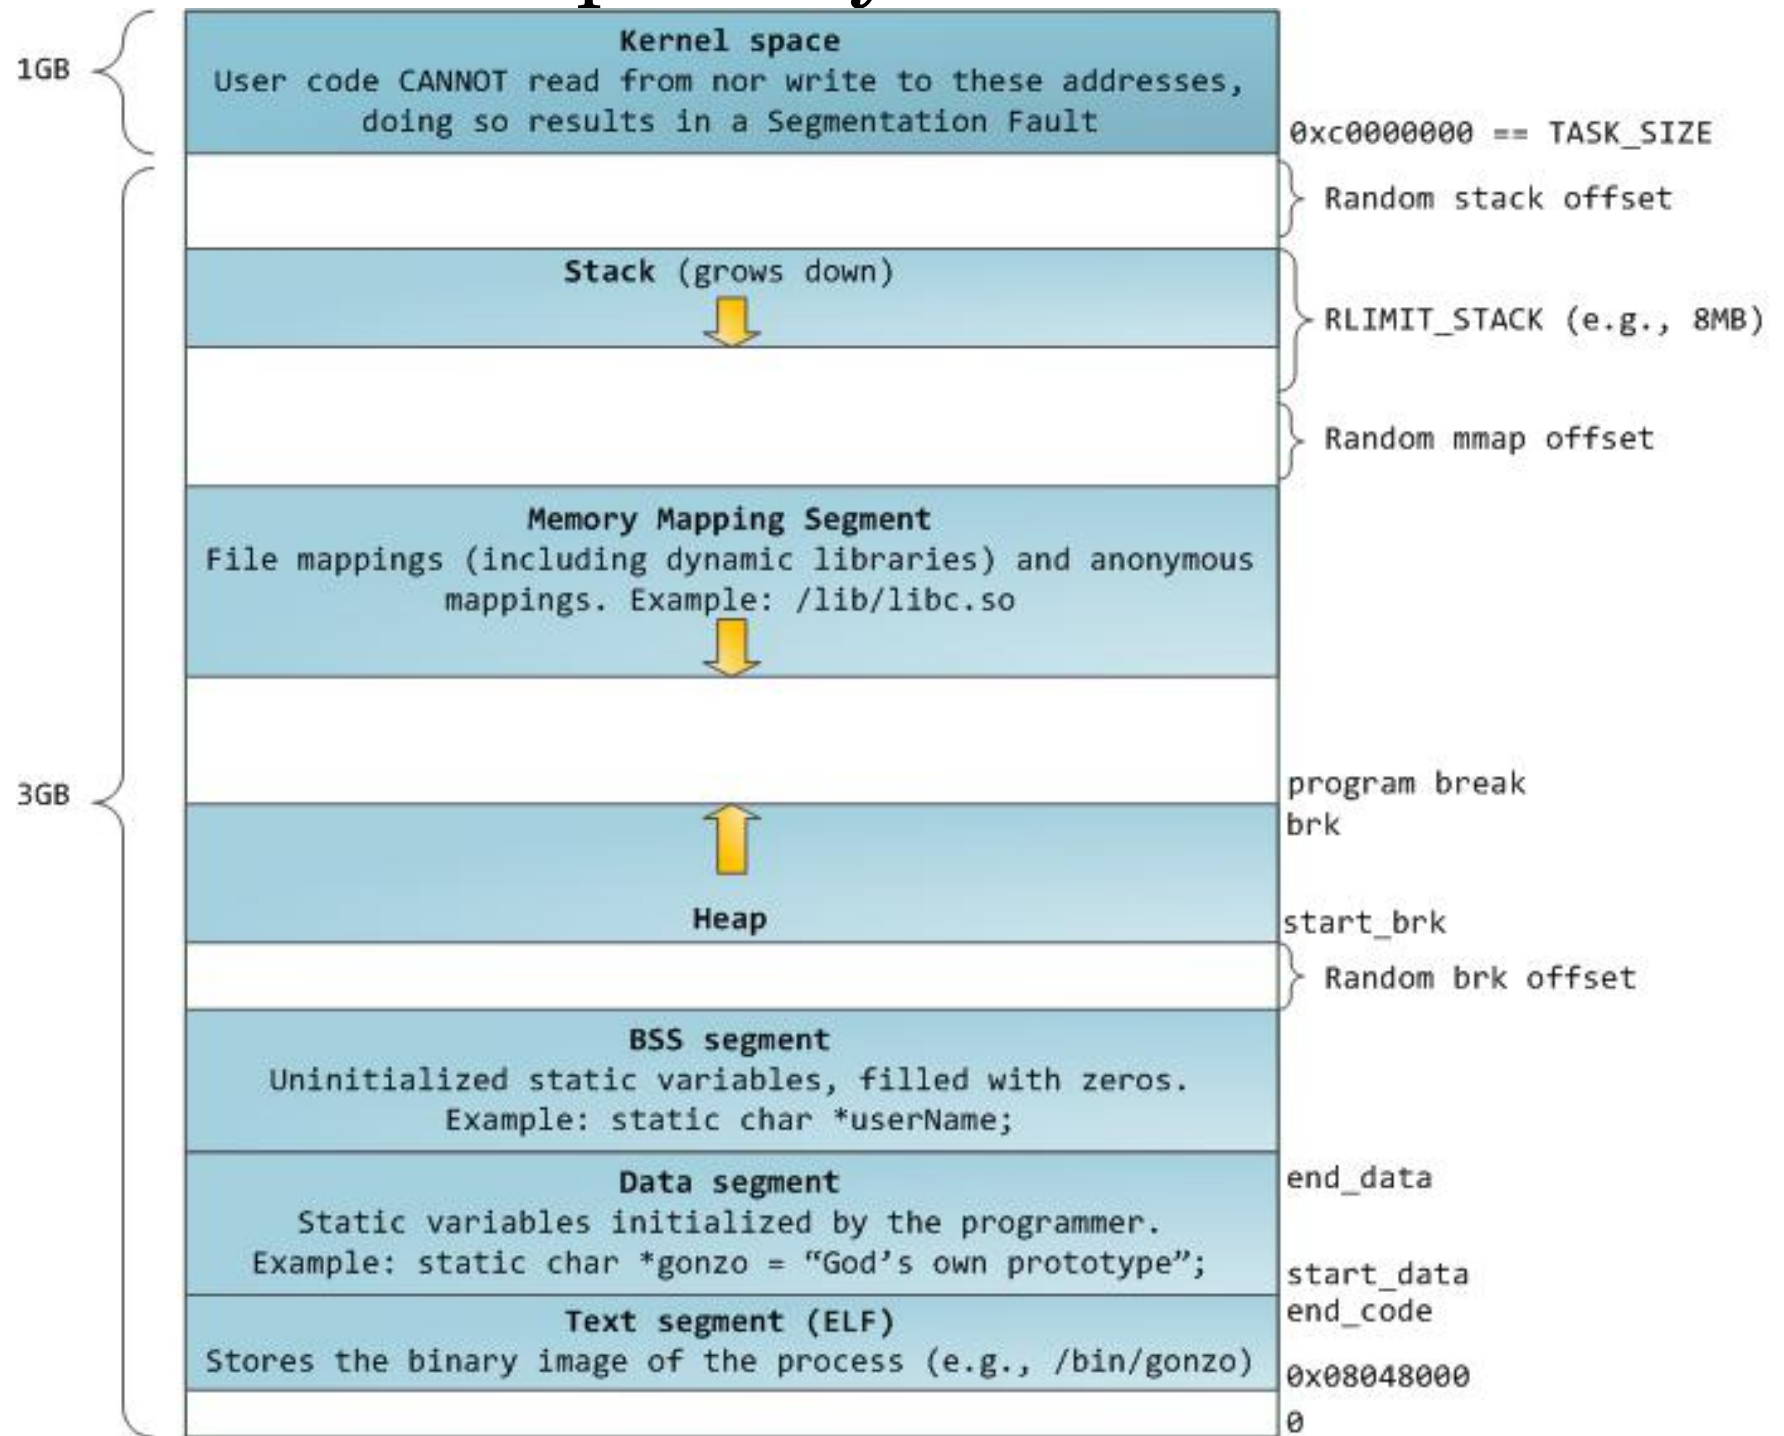


Code

00000000

Heap

# • A real address space layout



# C Memory API

- Malloc

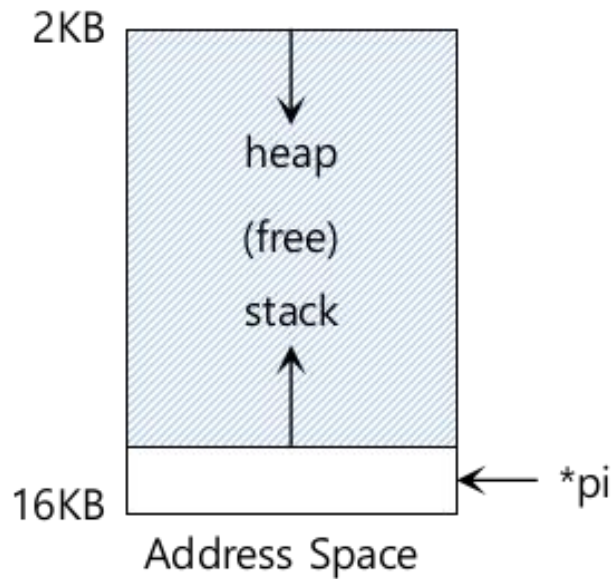
```
#include <stdlib.h>  
void *malloc(size_t size);
```

- If failed, return NULL

- Free

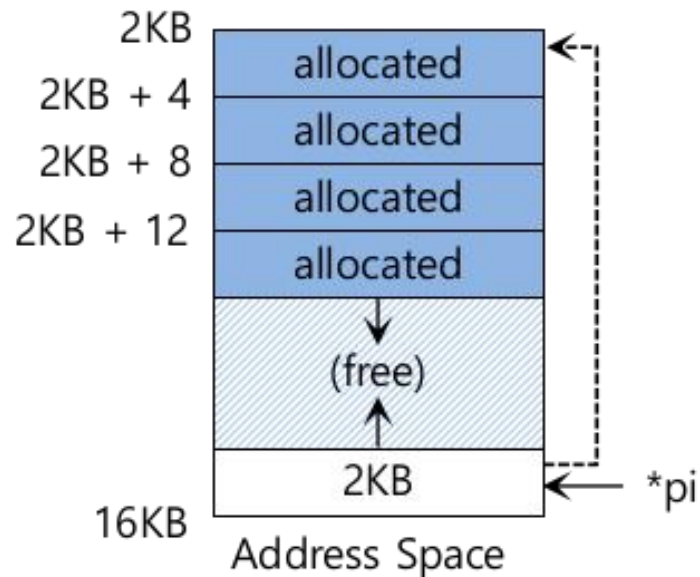
```
#include <stdlib.h>  
void free(void* ptr);
```

# Allocation



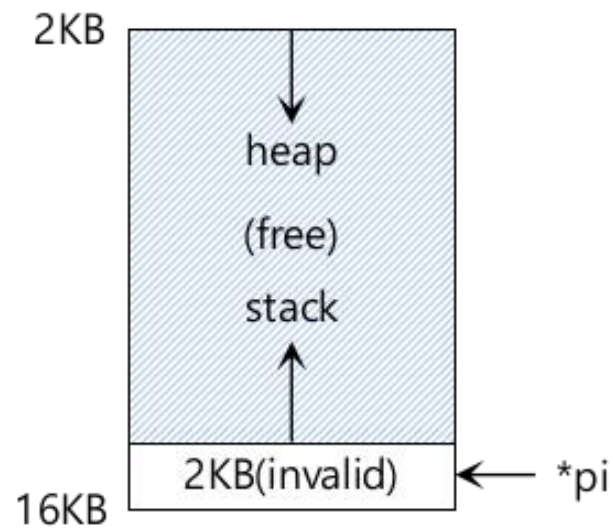
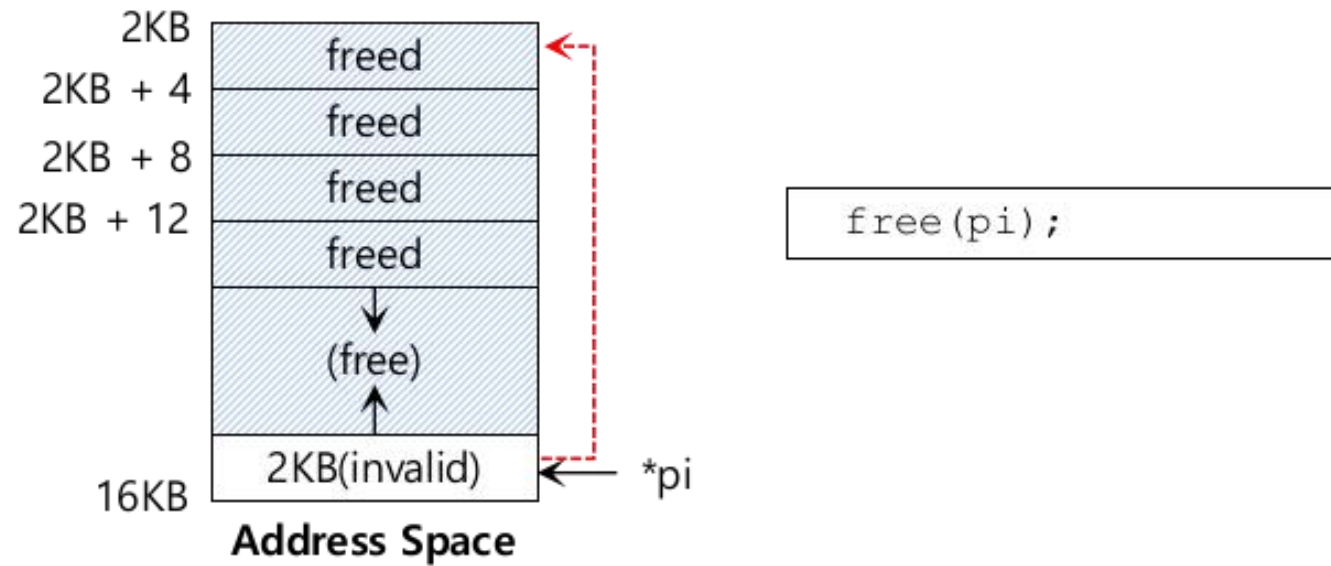
```
int *pi; // local variable
```

-----> pointer



```
pi = (int *)malloc(sizeof(int) *  
4);
```

# Free



# C Memory API

- Segment fault

```
char *src = "hello";  
char *dst;  
strcpy(dst, src);
```

- run this code, it will likely lead to a

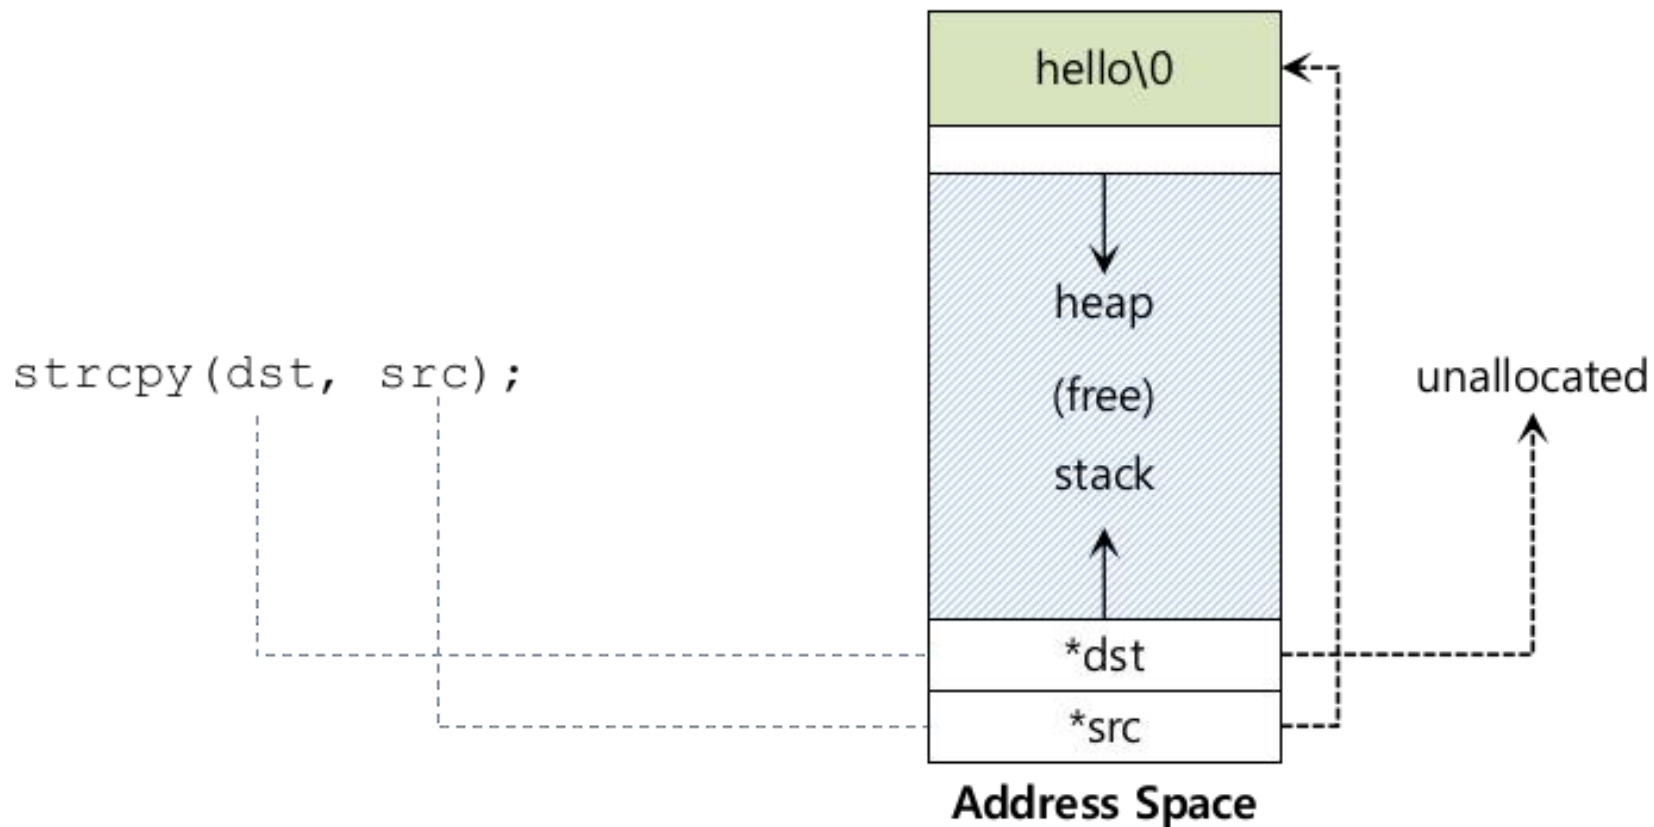
**segmentation fault**

- It is a fancy term for

YOU DID SOMETHING WRONG WITH MEMORY  
YOU FOOLISH PROGRAMMER AND I AM ANGRY.


# Segmentation Fault

```
char *src = "hello"; //character string constant  
char *dst;           //unallocated  
strcpy(dst, src);    //segfault and die
```



# Correct Code

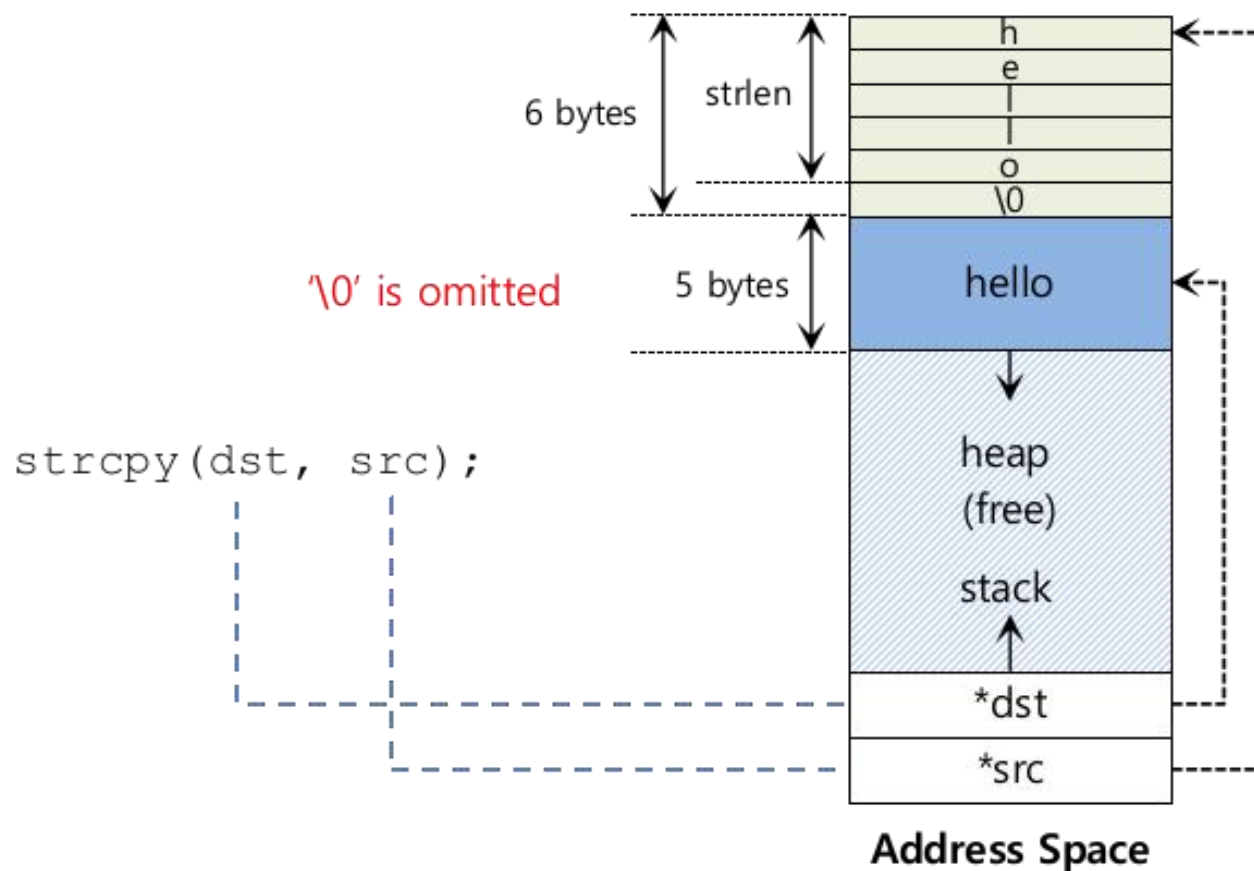
```
char *src = "hello";    //character string constant
char *dst (char *)malloc(strlen(src) + 1 ); // allocated
strcpy(dst, src);        //work properly
```





# Works, but buggy

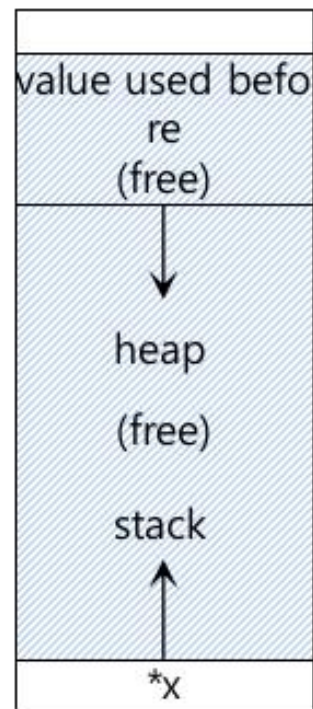
```
char *src = "hello"; //character string constant
char *dst (char *)malloc(strlen(src)); // too small
strcpy(dst, src);     //work properly
```



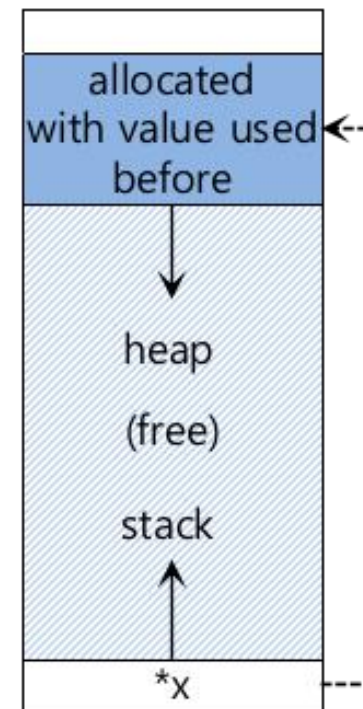
# Uninitialized Read

- Wild pointer

```
int *x = (int *)malloc(sizeof(int)); // allocated
printf("*x = %d\n", *x); // uninitialized memory access
```




Address Space




Address Space

# Memory Leak


**unused** : unused, but not freed



Address Space



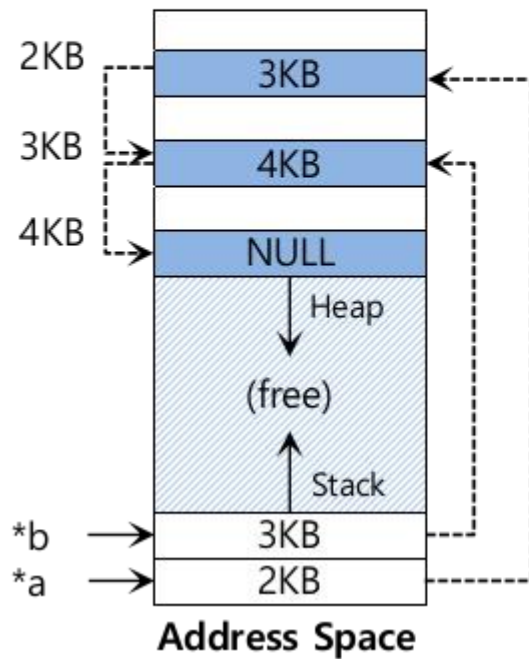
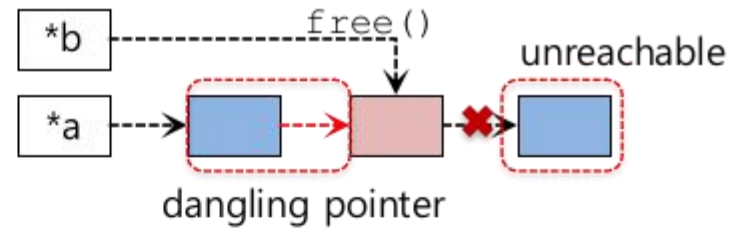
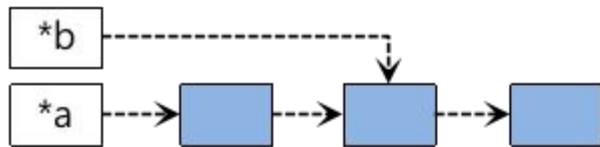
Address Space




Address Space

run out of memory

# Dangling Pointer



free  
(b)



# C Memory API

- Standard library
  - malloc(), realloc(), free()
- System calls
  - brk(), sbrk()
  - mmap()
- For comparison
  - printf() and write()
  - “Buffer the system call”

# System calls: brk(), sbrk()

```
#include <unistd.h>
```

```
int brk(void *addr)
```

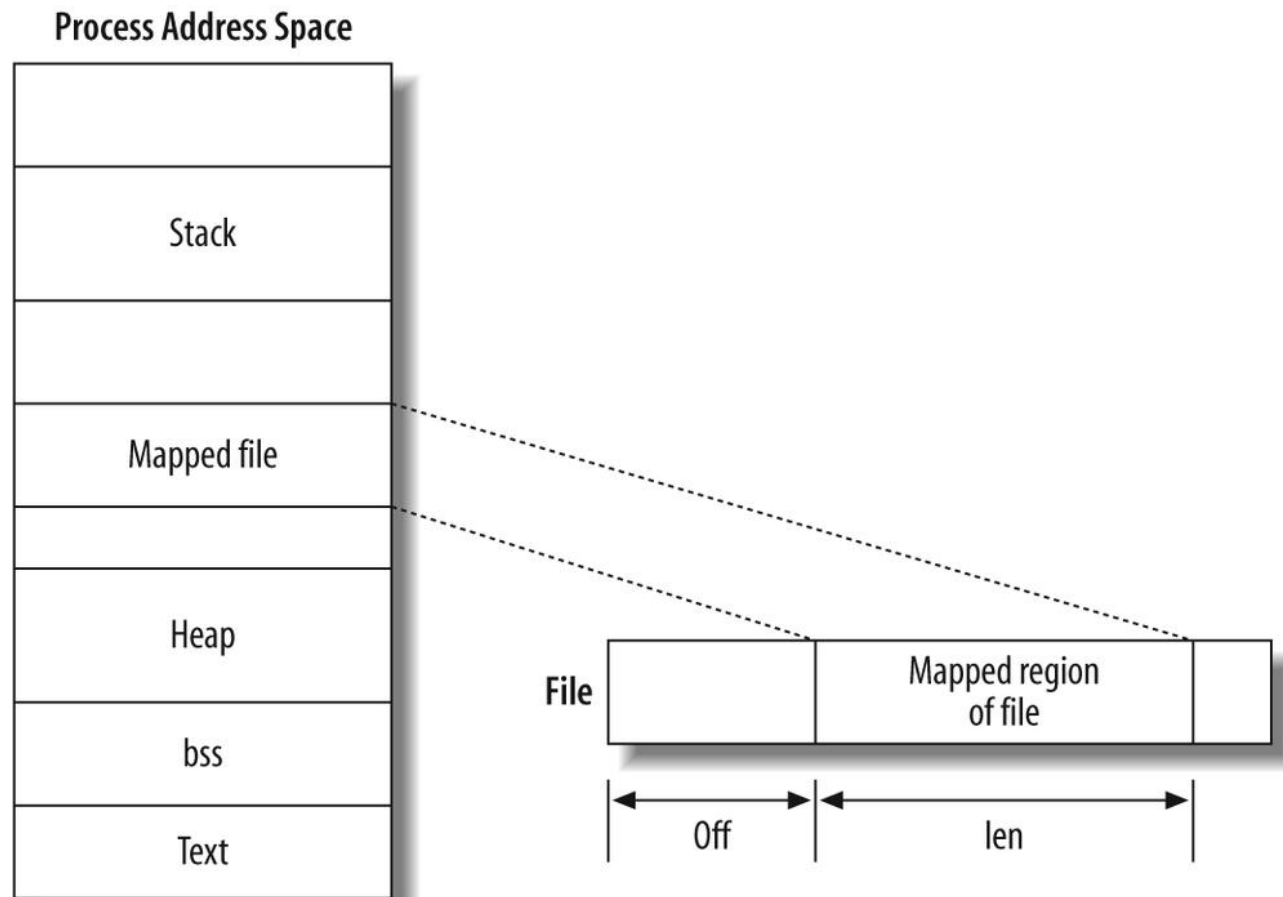
```
void *sbrk(intptr_t increment);
```

- brk/sbrk: expand the program's break.
  - break: The location of the end of the heap in address space

# System calls: mmap()

```
#include <sys/mman.h>
```

```
void *mmap(void *ptr, size_t length, int prot, int flags,  
int fd, off_t offset)
```




# C Memory API

- Summary: common errors
  - Forget to allocate memory
  - Not allocating enough memory
  - Forget to initialize allocated memory
  - Forget to free memory
  - Free memory before you are done with it
  - Free memory repeatedly
  - Call `free()` incorrectly





AND SUDDENLY YOU MISSTEP, STUMBLE, AND JOLT AWAKE?



# Free Memory Management



Dark Forest of Pointers

# Free Memory Management

- Manage a bunch of “free memory”
- Where ?
  - implementation of virtual memory (kernel space)
  - implementation of memory allocation librarys (user space)

# Free Memory Management

- Fixed-size unit
  - Paging
  - Problem: internal fragmentation
- Variable-size unit
  - User level: memory allocation library
  - Kernel level: VM implemented with segmentation
  - Problem: external fragmentation



# Free Memory Management

- Proj2 (linux part)
  - manage variable-size free memory units
  - implement a use space memory allocation library
    - `malloc(size_t size)`
    - `free(void *ptr)`

# Free Memory Management


- Assumptions
  - Focus on external fragmentation
  - No compaction
  - Manage a contiguous region of bytes (by `mmap()` system call)

# Free Memory Management

- Low-level Mechanisms
  - Splitting and Coalescing
  - Tracking allocated regions
  - Implementation of a free list
- High-level Intelligence
  - Best fit
  - Worst fit
  - First fit
  - Next fit

# Free Memory Management

- **Splitting and Coalescing**
  - Free list: a set of free chunks
  - Two chunks (10 bytes each)







# Free Memory Management

- **Splitting and Coalescing**

- request less than 10 bytes? (e.g. malloc(1))



- **Splitting**



# Free Memory Management


- **Splitting and Coalescing**

- Free a chunk?



- Malloc(20)?

- **Coalescing**



# Free Memory Management

- **Tracking Allocated Regions**

- Observation on `free(void *ptr)`
  - No size parameter
- Given a pointer, the malloc library could determine the size of region
- How?
  - Some extra information
  - **header** of a memory block


# Free Memory Management

- **Tracking Allocated Regions**

- header


```
typedef struct __header_t {  
    int size;  
    int magic;  
} header_t;
```

- malloc(20)



# Free Memory Management

- **Tracking Allocated Regions**
  - header: example



# Free Memory Management

- **Tracking Allocated Regions**

- free(ptr)

- Get the size of the region

```
void free(void *ptr) {  
    header_t *hptr = (void *)ptr - sizeof(header_t);  
}
```


- Check whether ptr is valid

```
assert(hptr->magic == 1234567)
```

# Free Memory Management

- **Implementation of the Free List**

- Free list



- Implementation

- List node (allocate a node when needed)
    - Can NOT do this here! All you have is a given free space
  - How to build a free list **inside** the free space?

# Free Memory Management

- **Implementation of the Free List**
  - Node in free list


```
typedef struct __node_t {  
    int size;  
    struct __node_t *next;  
} node_t;
```



# Free Memory Management

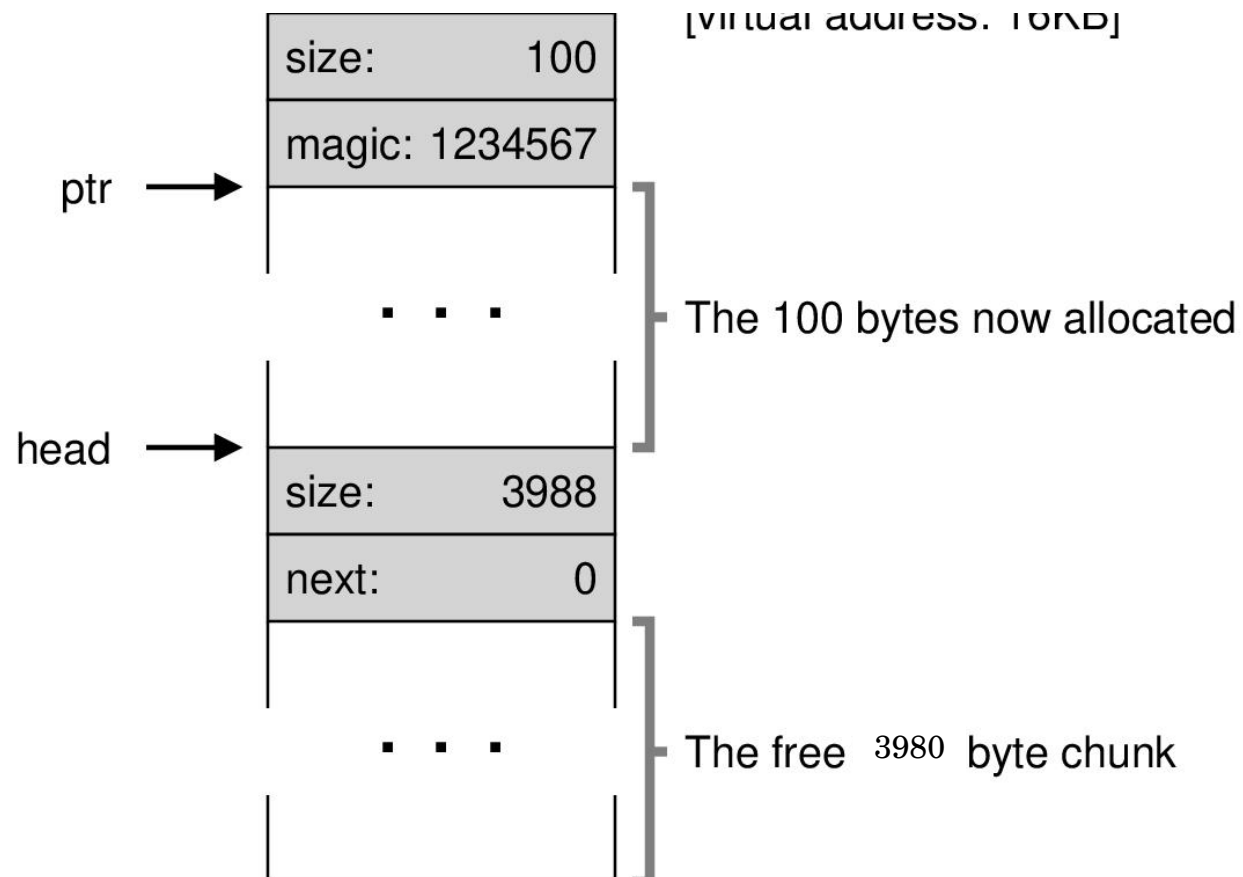
- **Implementation of the Free List**
  - Initialization (e.g. 4096)

```
// mmap() returns a pointer to a chunk of free space
node_t *head = mmap(NULL, 4096, PROT_READ |
                    PROT_WRITE, MAP_ANON | MAP_PRIVATE, -1, 0);
head->size = 4096 - sizeof(node_t);
head->next = NULL;
```

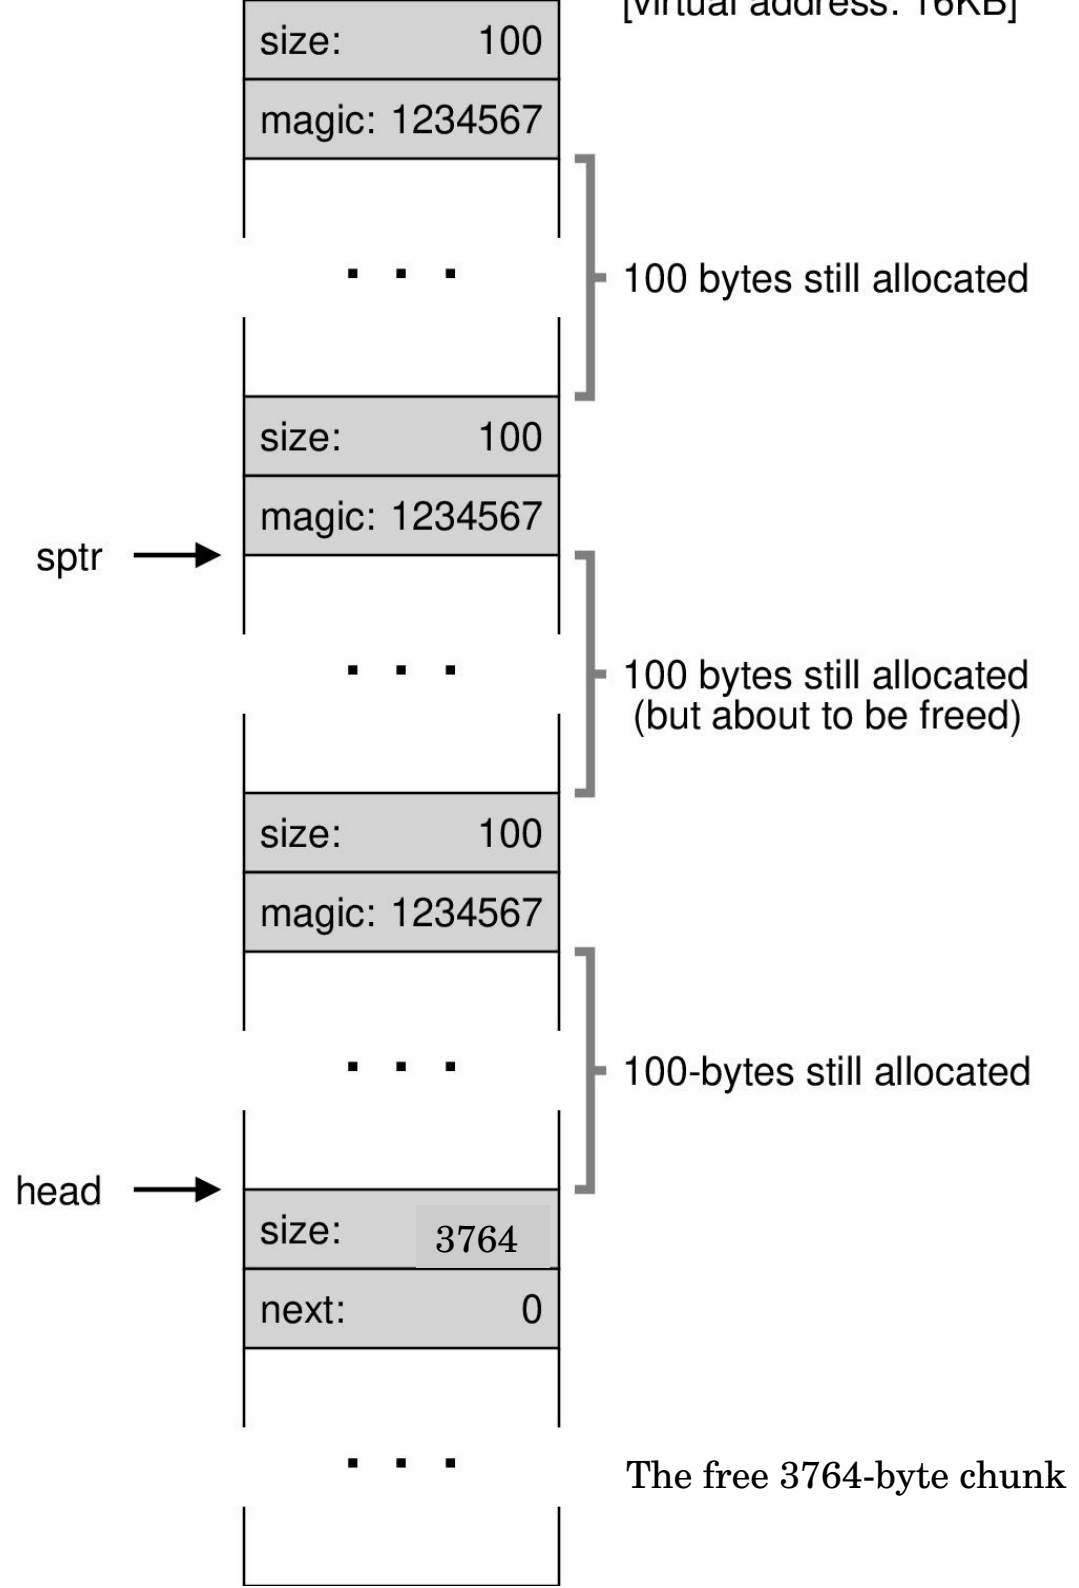


# Free Memory Management

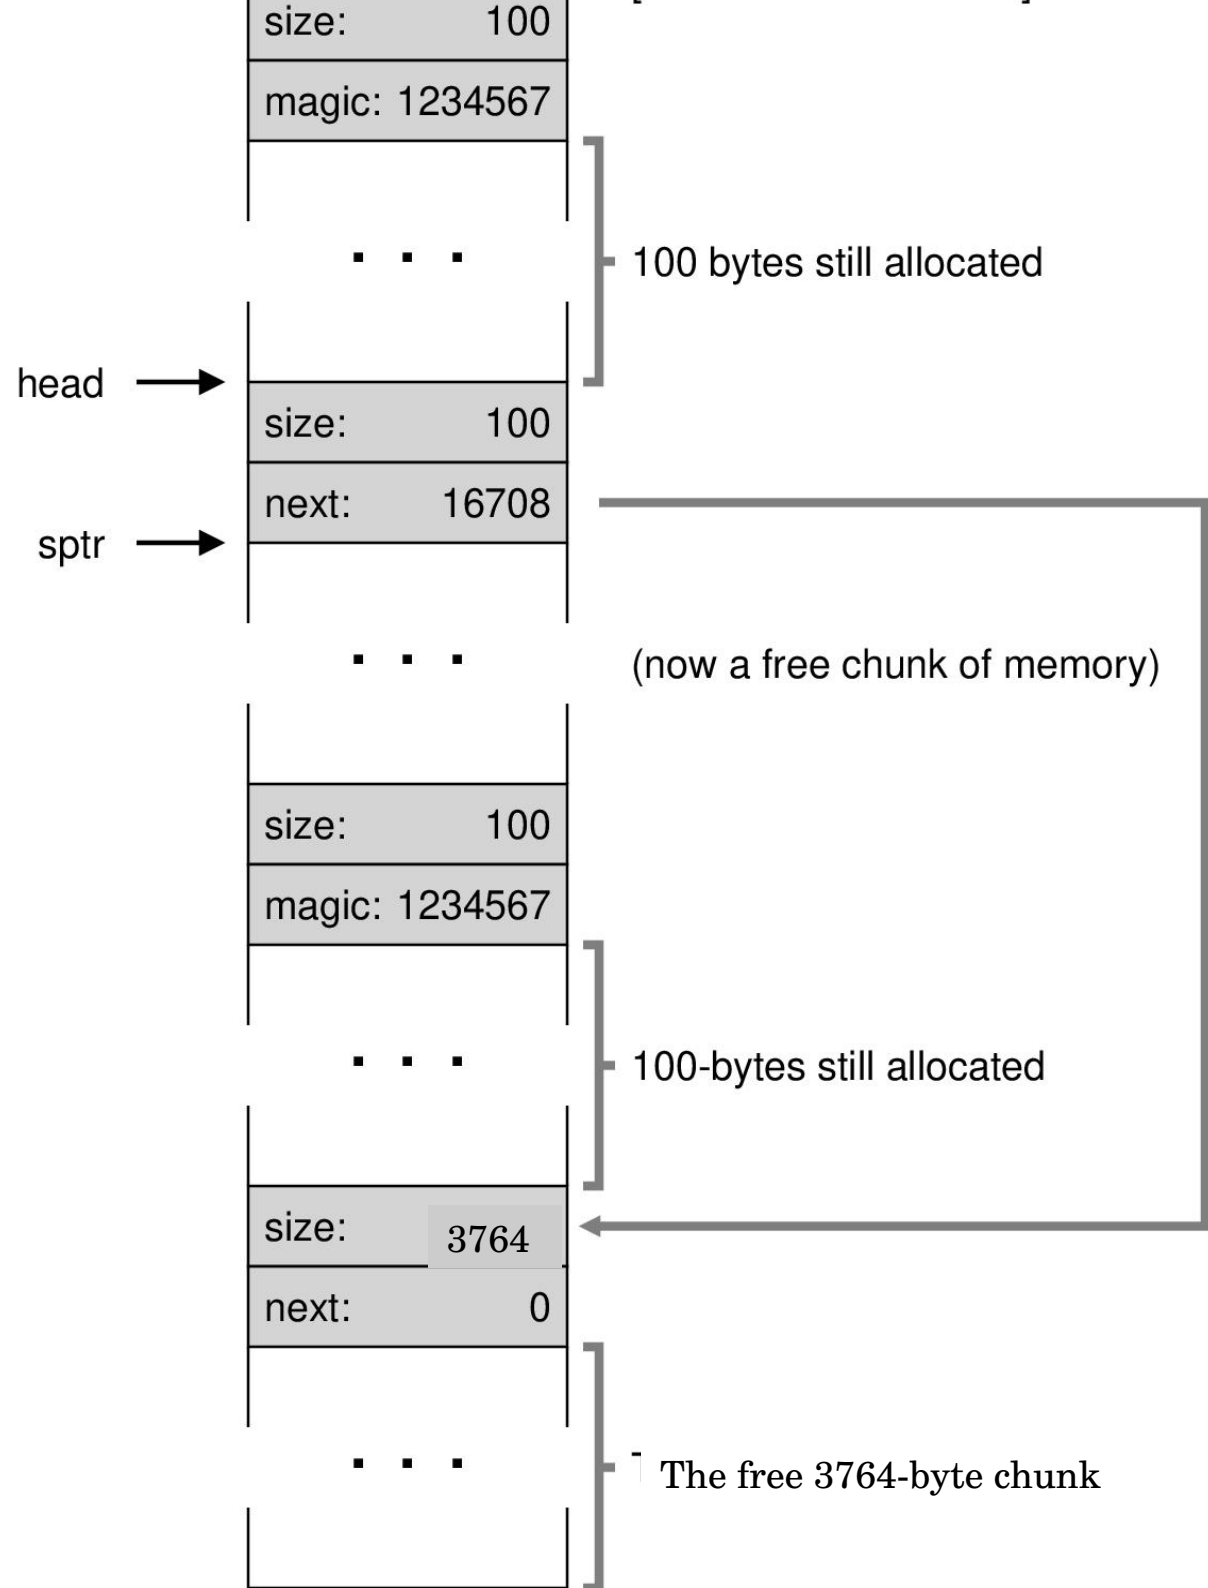
- **Implementation of the Free List**
  - `malloc(100)`



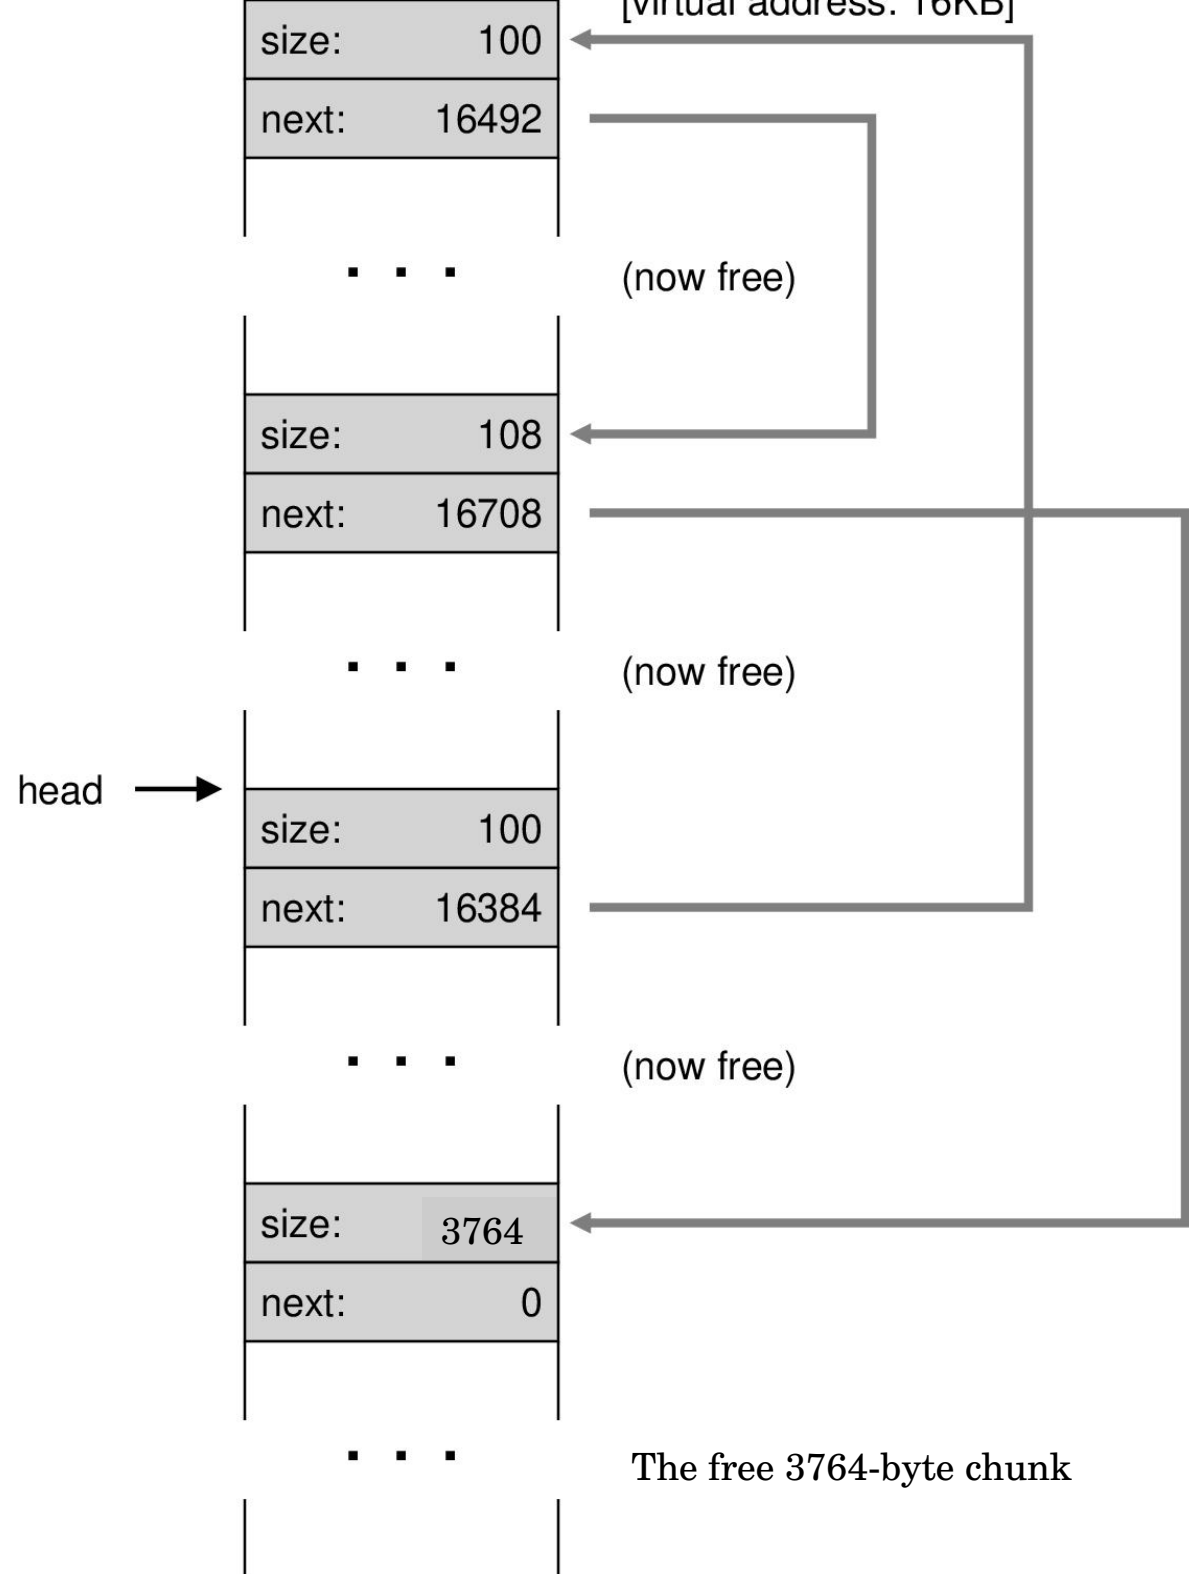
- `malloc(100)*3`



- **Free(16500)**  
 –  $16384 + 108 + 8$



- Free()\*3
- Coalesce
  - Merge adjacent
  - chunks



# Free Memory Management

- Growing the Heap
  - What if the heap runs out of space?
    - Return NULL
  - Increase the size of heap
    - OS find free physical pages
    - Map them into address space of the process

# Free Memory Management

- Summary of low-level Mechanisms
  - Splitting and Coalescing
  - Tracking allocated regions
  - Implementation of a free list
  - Growing the heap

# Free Memory Management

- High-level intelligence
  - How to find the proper nodes in the free list?
    - Less fragmentation
    - Fast allocation
  - Some simple strategies
    - The stream of allocation and free requests can be arbitrary
    - Any strategy could be arbitrarily bad/good



# Free Memory Management

- Best Fit
  - Find the smallest feasible node
- Worst Fit
  - Find the largest feasible node
- First Fit
  - Find the first feasible node

# Free Memory Management

- Example



- Best fit




- Worst fit



# Free Memory Management

- Other approaches
  - Segregated List
    - Slab allocator
  - Buddy Allocation
    - Binary search tree



# Free Memory Management

- Dlmalloc (Doug Lea allocator)
  - Segregated list for small size allocations
  - Search the free list
  - sbrk and mmap
  - ...
- <http://g.oswego.edu/dl/html/malloc.html>
- [https://cs61.seas.harvard.edu/wiki/images/e/e2/Lec11-Dynamic\\_memory\\_2.pdf](https://cs61.seas.harvard.edu/wiki/images/e/e2/Lec11-Dynamic_memory_2.pdf)