

Operating System Labs

Yuanbin Wu
cs@ecnu

Oral Tests

- Project 2 oral test
 - Date: Nov. 27
- How
 - 5min presentation
 - 2min Q&A

Oral Tests

- Who
 - Principle: you should take at least one oral test
 - We assume that you know all design/implementation details about your project

Oral Tests

- Examples organization of presentations
 - What have you done?
 - Project background
 - How did you accomplish them?
 - data structures, algorithms,
 - Your favorite parts.
 - Features that you've tried, but failed
 - What did you learn from the project?
 - Possible future improvements
 - ...
- Highlight your new features (of part a or part b)

Oral Tests


- Suggestions for your slides
 - The clew model and onion model
 - Minimize words, maximize pictures
 - Simple and clear
 - Large font
- Suggestions for your talk
 - If you are an audience of your own talk...
 - Design your rhythm, pauses, actions...
 - Practice

Oral Tests

- How to speak, by Patrick Winston (MIT)
 - <https://www.bilibili.com/video/BV1K54y1m7M6?from=search&seid=3207037721399021621>
- Suggestions from Jonathan Shewchuk (UC Berkeley)
 - <http://www.cs.berkeley.edu/~jrs/speaking.html>

Operating System Labs

“I am trained to only sleep during national holidays”




Operating System Labs

- Review of Memory Management

Memory Management

- Early days




Memory Management

- Multiprogramming
 - multiple processes could be ready to run at a given time
 - the OS would switch between them
- Time sharing
 - many users might be concurrently using a machine

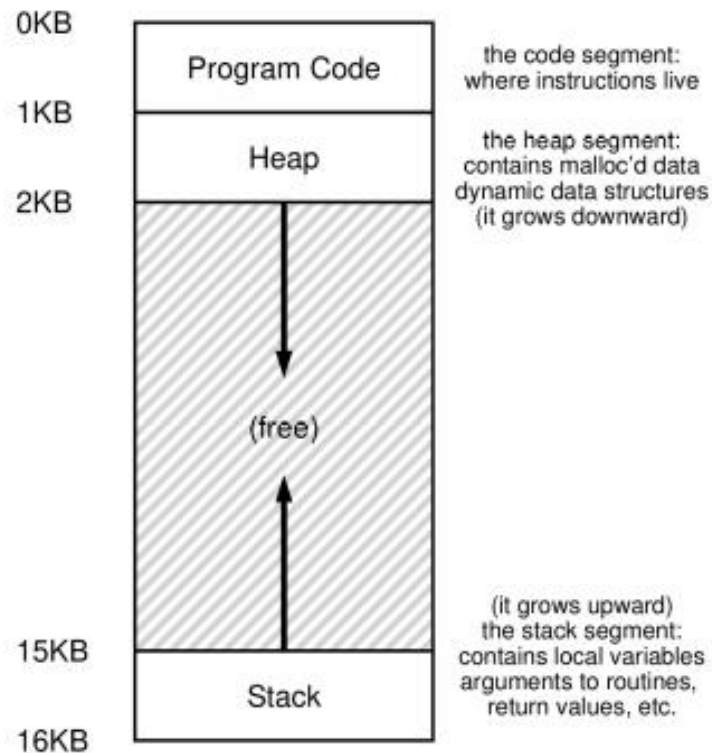
Memory Management

- Multiprogramming and Time Sharing
 - Multiple processes live in memory simultaneously





Memory Management

- Multiprogramming requires easy-to-use virtualization of memory
 - A concept called “address space”



Memory Management

- Two views on memory
 - From processes: different processes have different address spaces
 - From OS: limited physical memory cells



Memory Management

- Memory management
 - How OS provides such easy-to-use address spaces for processes?
 - **Virtualization** of memory
 - Recall: virtualization of CPU

Memory Management

- Goals of Virtualize Memory(and other things)
 - Transparency
 - Efficiency
 - Protection
 - The OS should make sure to protect processes from one another

Memory Management

- Transparency
 - OS should implement virtual memory in a way that is invisible to the running program
 - From the programmer's point of view:
 - Every address is fraud
 - Only the OS knows the truth

Memory Management

- Virtualize Memory: **Limited Direct Execute**
 - **Hardware-based address translation**
 - Hardware:
 - transparency, efficiency, protection
 - OS:
 - configure hardware correctly
 - manage free memory
 - handle exception

Memory Management


- Hardware: **Transparency**
 - We starts with a simple idea called
 - Base and bounds
 - Dynamical (hardware-based) allocation

An Example


```
void func ()  
{  
    int x;  
    x = x + 3;  
}
```

```
128: movl 0x0(%ebx), %eax    ;load 0+ebx into eax  
132: addl $0x03, %eax        ;add 3 to eax register  
135: movl %eax, 0x0(%ebx)    ;store eax back to mem
```

Fetch instruction at address 128
Execute this instruction (load from address 15 KB)
Fetch instruction at address 132
Execute this instruction (no memory reference)
Fetch the instruction at address 135
Execute this instruction (store to address 15 KB)



Address space



Hardware:


- 2 registers in CPU
- **Base**: the start of phy mem
- **Bound**: the size of phy mem



$$\text{physical} = \text{virtual} + \text{base}$$

Base: 32K
Bound: 16K

Physical Memory



physical = virtual + base

Fetch instruction at address 128

Execute (load from address 15 KB)

Fetch instruction at address 132

Execute (no memory reference)

Fetch the instruction at address 135

Execute (store to address 15 KB)

Visiting address 128

$$\begin{aligned} &128 + 32K \\ &= 128 + 32768 \\ &= 32896 \end{aligned}$$


Base: 32K

Bound: 16K

Address Space



Physical Memory



physical = virtual + base

Fetch instruction at address 128
Execute (load from address 15 KB)
Fetch instruction at address 132
Execute (no memory reference)
Fetch the instruction at address 135
Execute (store to address 15 KB)

128: movl 0x0(%ebx), %eax


$$\begin{aligned} &15K + 32K \\ &= 47K \end{aligned}$$

Base: 32K
Bound: 16K

Address Space



Physical Memory



Memory Management

- Hardware: **Protection**
 - Bounds reg
 - Raise an exception when the required address is illegal
 - Know how to do when exceptions are raised
 - E.g.
 - Base: 0
Bound: 4K
 - Then address 4400 is illegal according to the Bound

Memory Management

- Hardware: **Efficiency**
 - The registers are in CPU chip
 - The part of CPU related to address translation is called: **MMU** (memory management unit)

Memory Management

- Hardware requirements summary
 - Privileged mode
 - Base/bounds registers
 - Ability to translate virtual addresses and check if within bounds
 - Privileged instruction(s) to update base/bounds
 - Privileged instruction(s) to register exception handlers
 - Ability to raise exceptions

Memory Management

- OS:
 - Maintain a data structure: **free list**
 - Find place in physical memory for a process when creating it
 - Collect the space when a process terminate
 - Context switch
 - Correctly configure base / bound register
 - Handle exception

**OS @ boot
(kernel mode)**

Hardware

initialize trap table

remember addresses of...
system call handler
timer handler
illegal mem-access handler
illegal instruction handler

start interrupt timer

start timer; interrupt after X ms

initialize process table

initialize free list

OS @ run
(kernel mode)

Hardware

Program
(user mode)

To start process A:

allocate entry in process table

allocate memory for process

set base/bounds registers

return-from-trap (into A)

restore registers of A

move to **user mode**

jump to A's (initial) PC

Process A runs

Fetch instruction


...

Memory Management

- Two implementation of virtual memory
 - Segmentation
 - Paging

Segmentation

- The problem of Base and Bound
 - Load entire address space
 - Wasteful
 - How to support large address space




Segmentation

- Solution:
 - Multiple base/bound
 - 3 logical segmentations
 - Code
 - Stack
 - Heap
 - 3 groups of base/bound registers

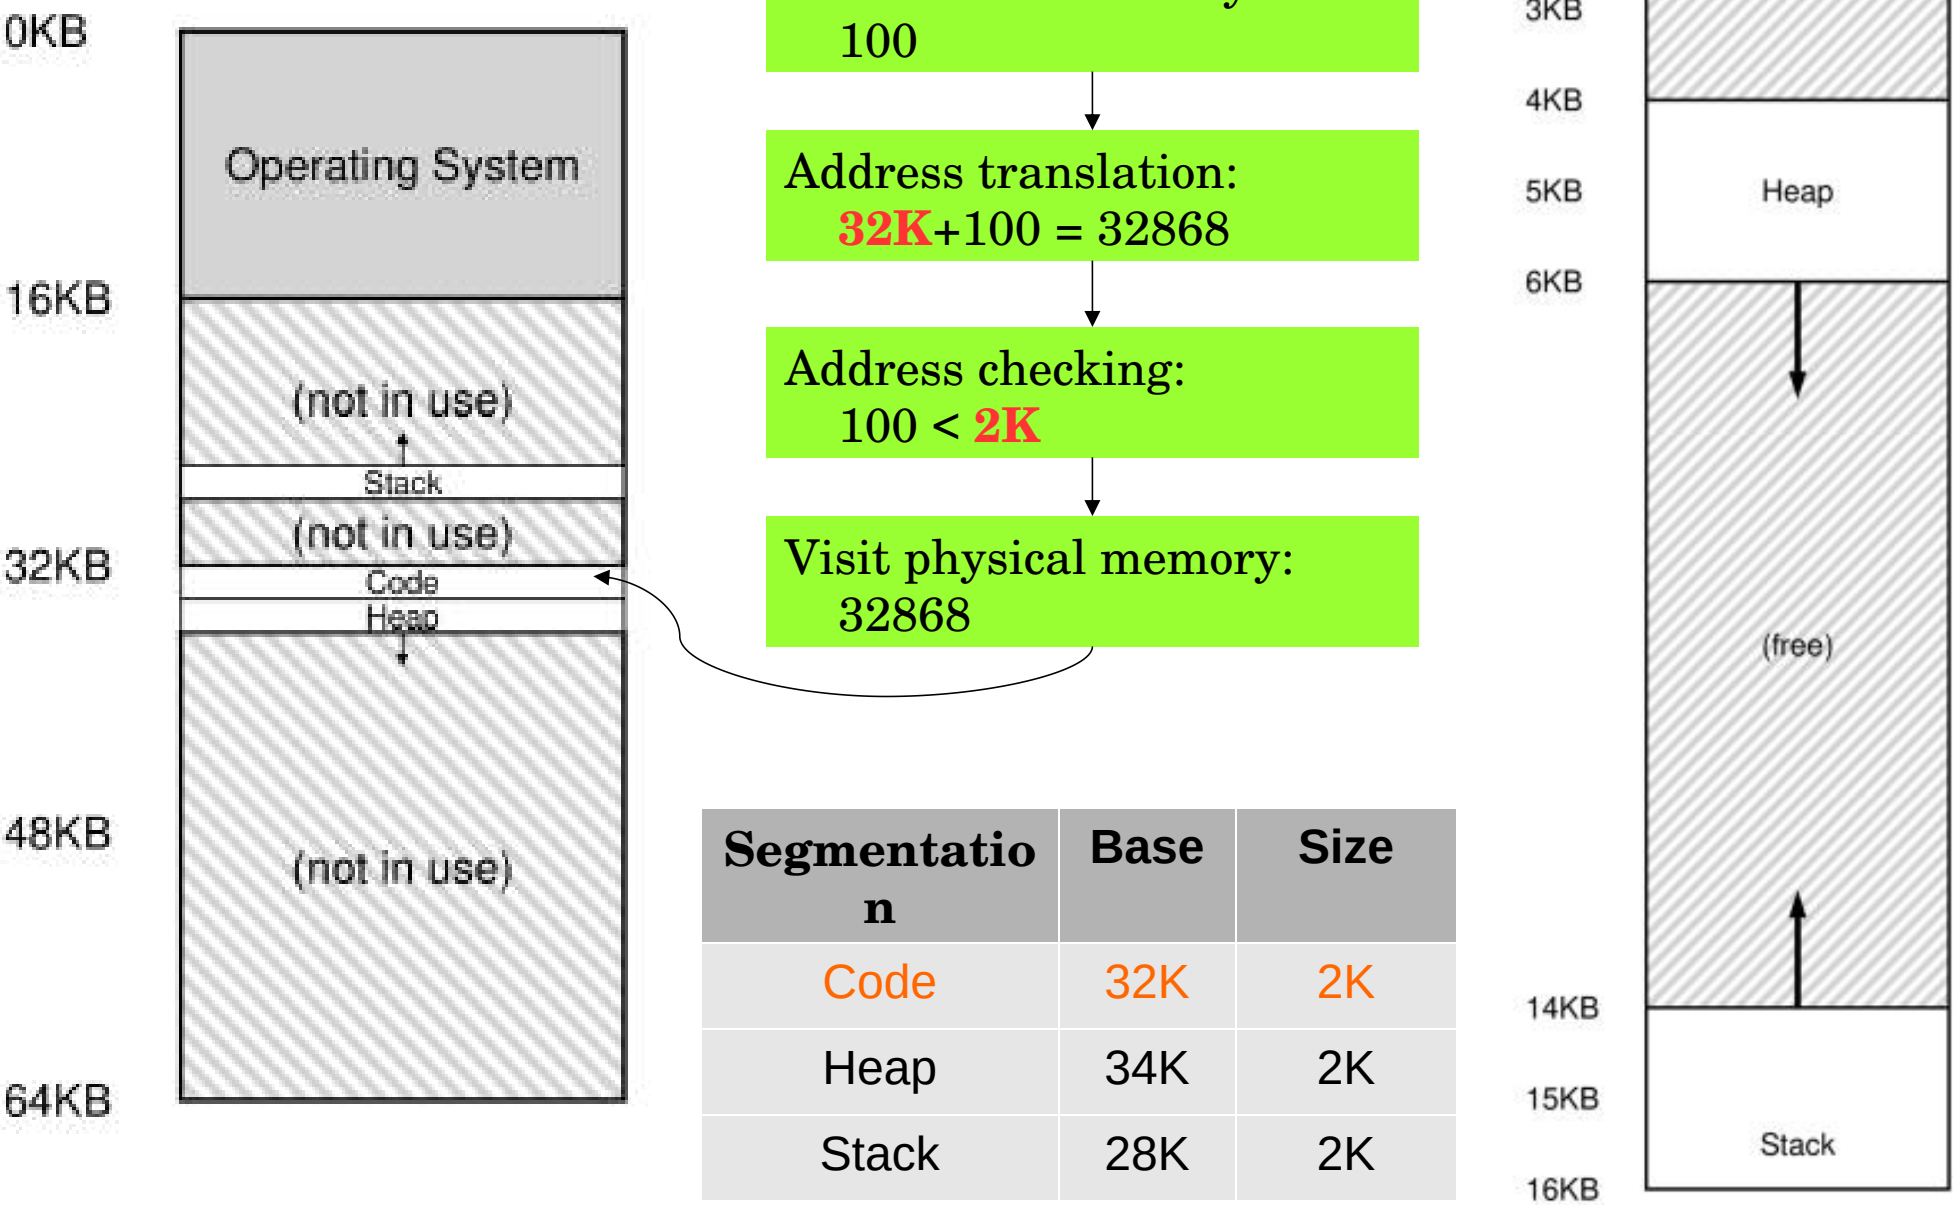
Segmentation

- Multiple base/bound
 - Physical memory

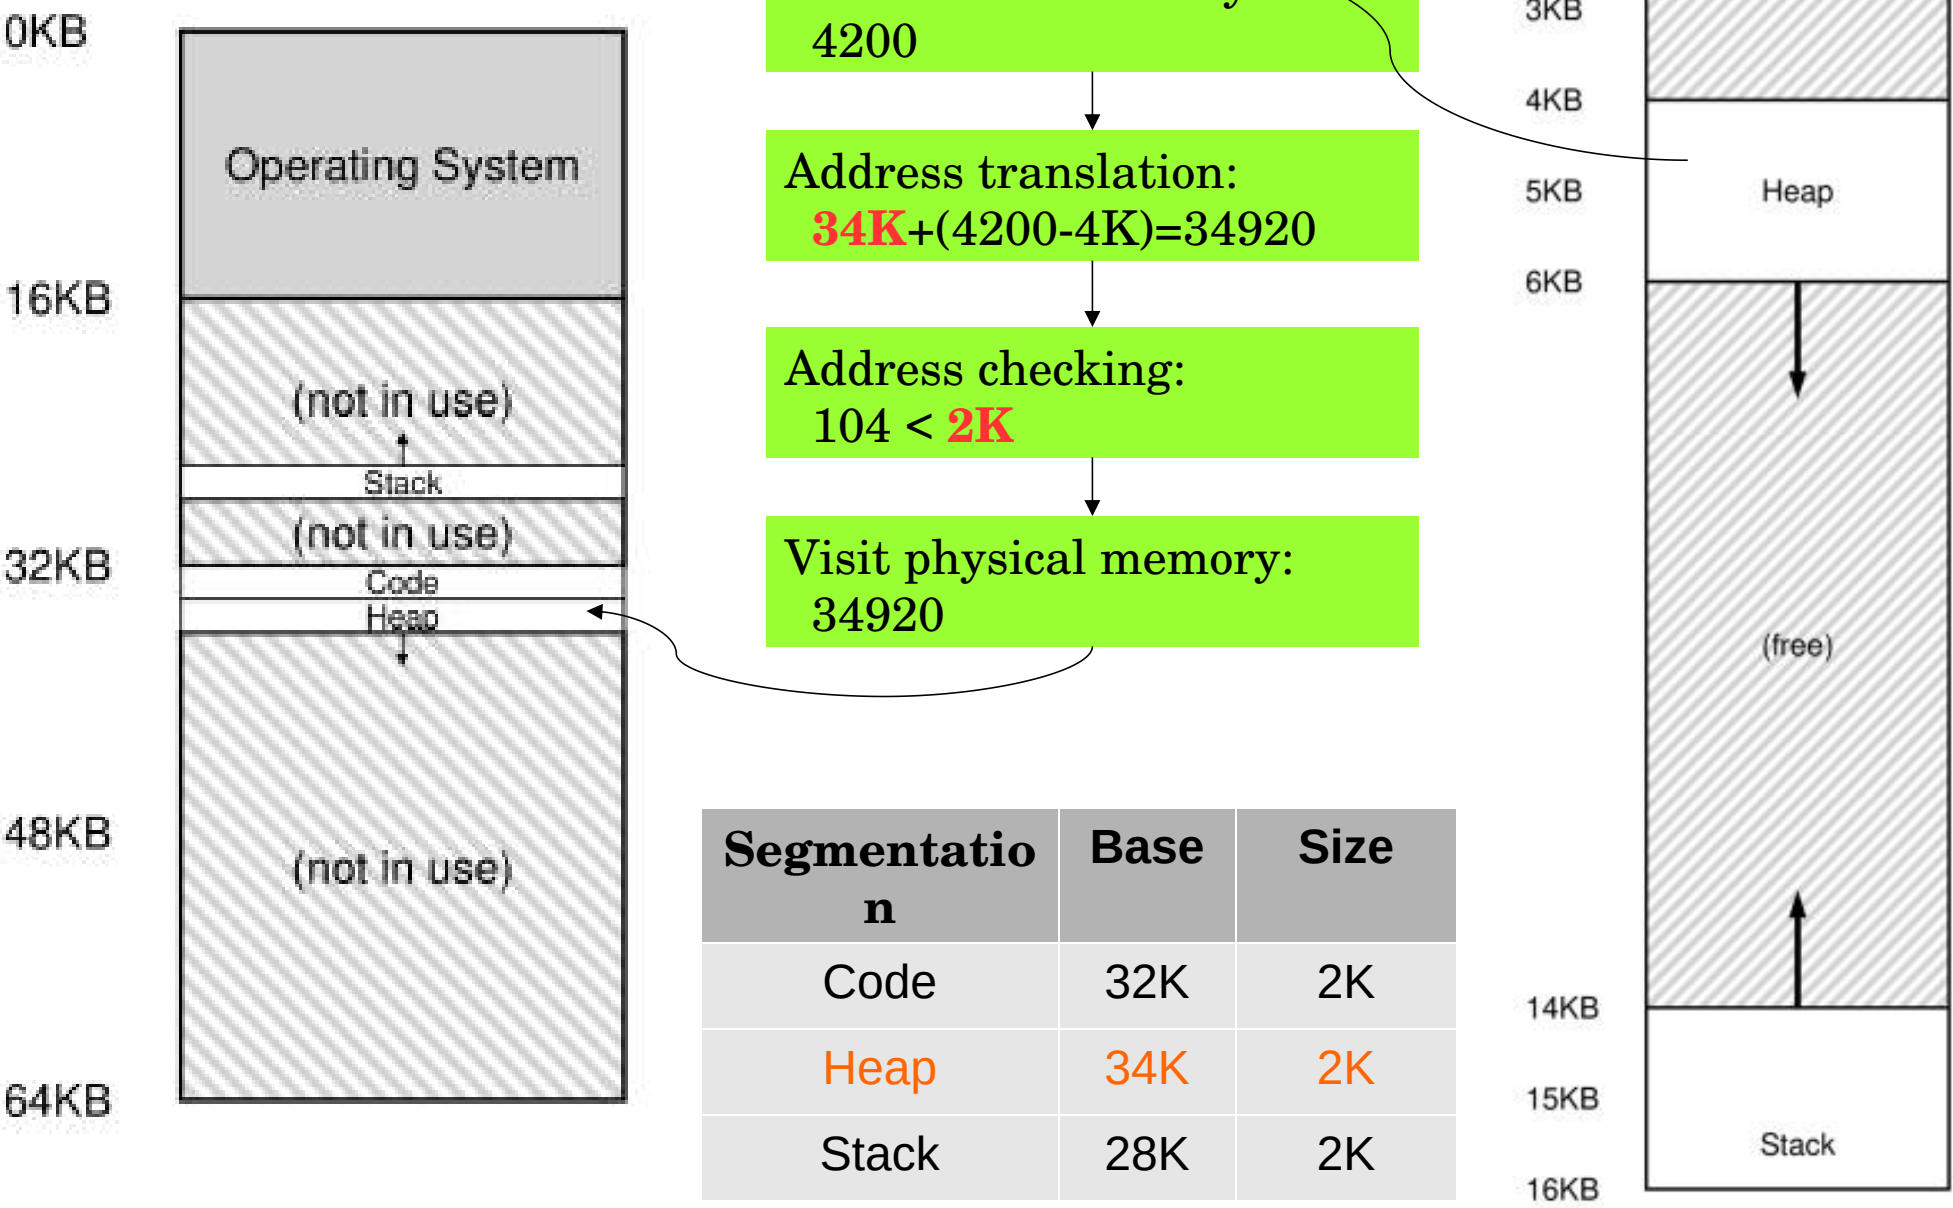
Segmentation	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K



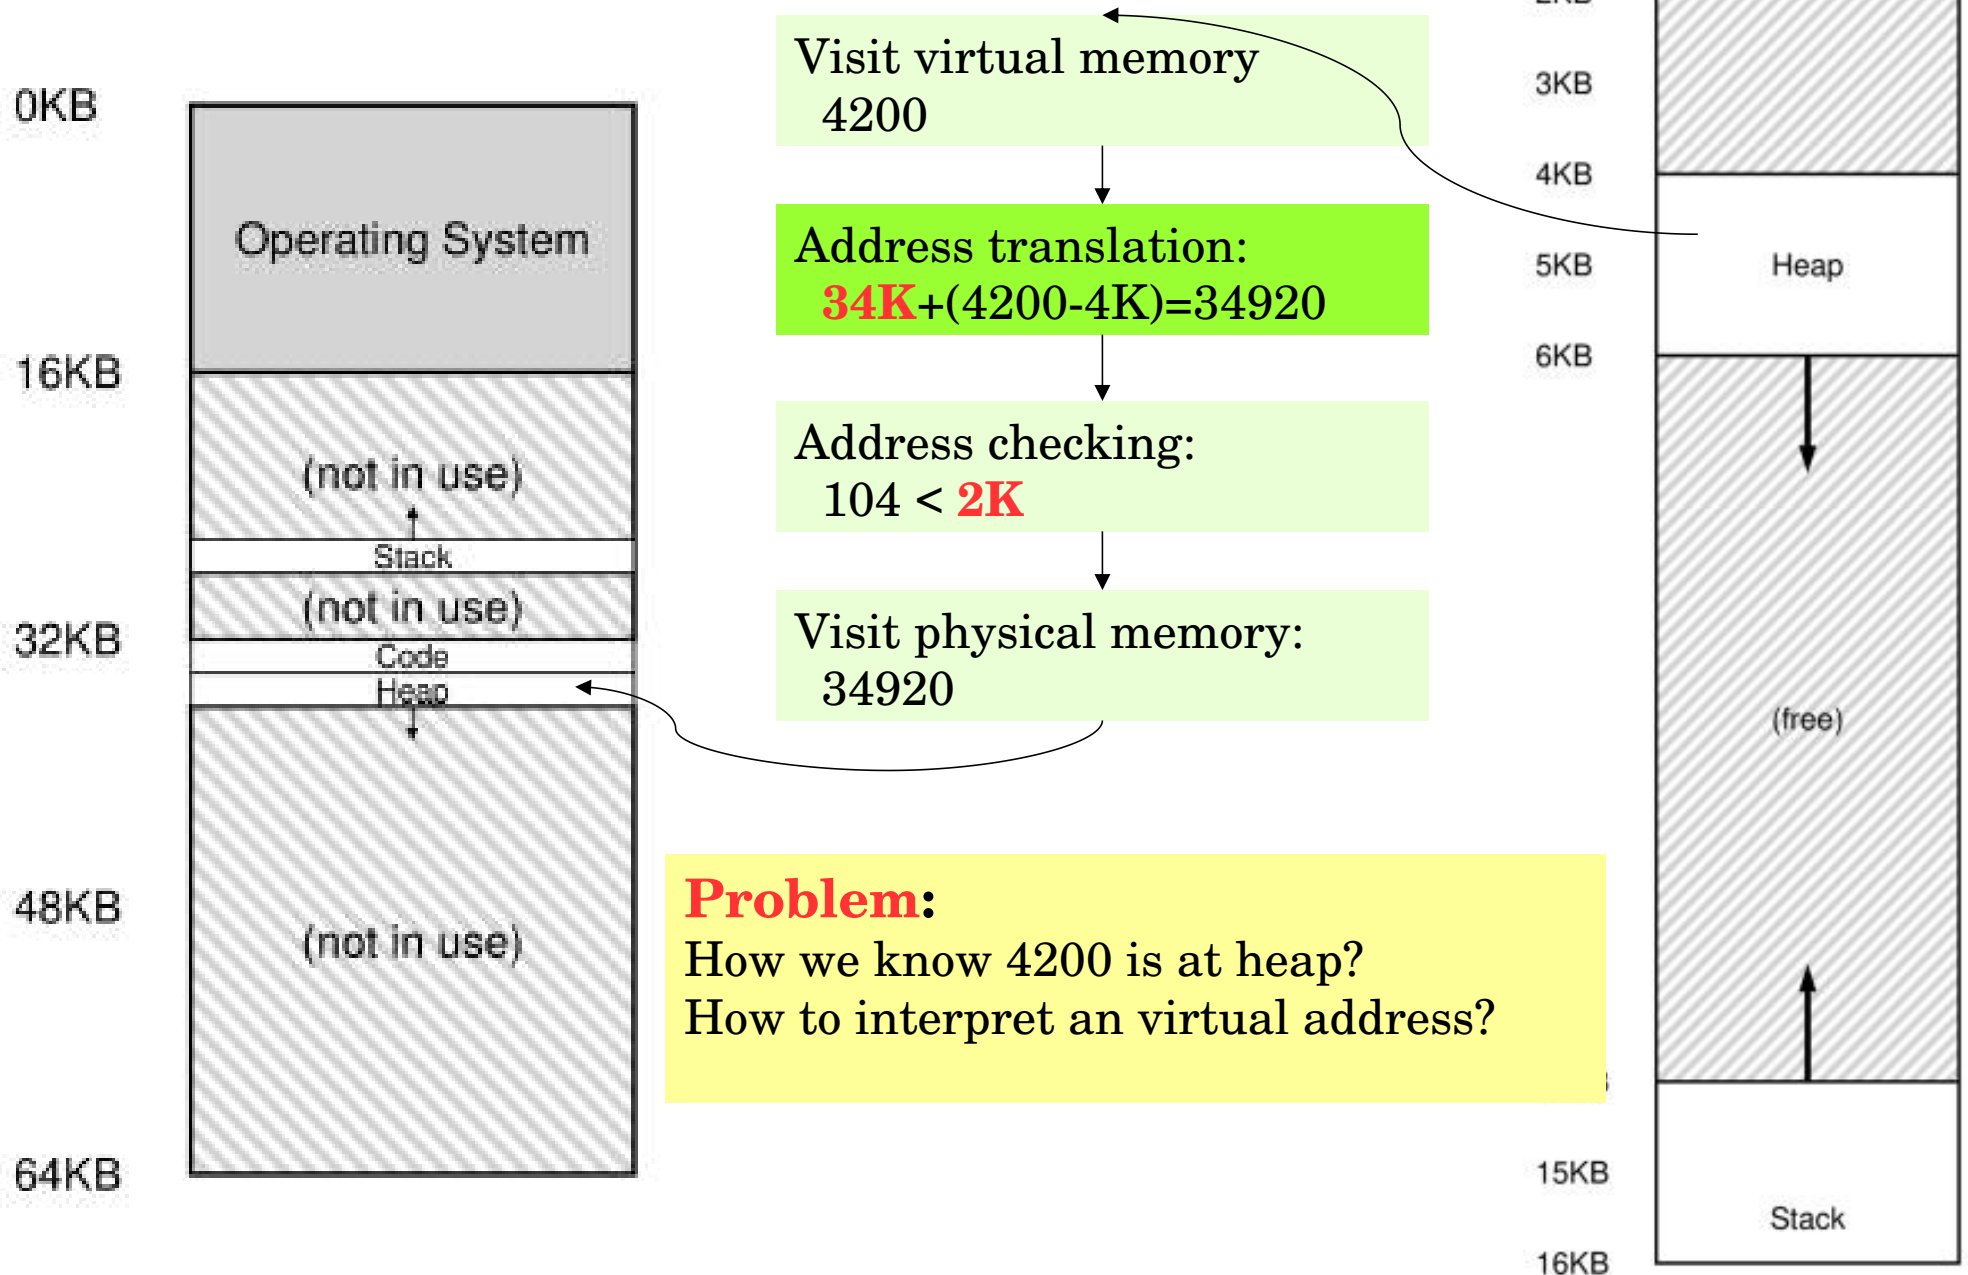
Example: multiple base/bound



Example: multiple base/bound



Example: multiple base/bound



Segmentation

- Which segmentation are we referring to
 - Explicit approach
 - top few bits of the virtual address
 - Example:
 - 16K address space \rightarrow 14 bit

Segmentation

- Which segmentation are we referring to
 - Example: 4200

Segmentation

- Which segmentation are we referring to

```
// get top 2 bits of 14-bit VA
Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT


// now get offset
Offset = VirtualAddress & OFFSET_MASK
if (Offset >= Bounds[Segment])
    RaiseException(PROTECTION_FAULT)
else
    PhysAddr = Base[Segment] + Offset

Register = AccessMemory(PhysAddr)
```

Segmentation

- About the stack
 - Difference
 - growth backwards
 - 28K - 26K

Segmentation	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K



Segmentation

- About the stack
 - Solution: extra hardware support
 - one bit in MMU
 - 1: growth in positive direction
 - 0: growth in negative direction

Segmentation	Base	Size	Grows Postive
Code	32K	2K	1
Heap	34K	2K	1
Stack	28K	2K	0

Example: multiple base/bound

Visit virtual memory

15K **11** **11** 00 00 00 00 00

Address translation:

[1] segment = 11 → stack reg

[2] offset = 3K

[3] maximum segment = 4K

[4] $3K - 4K = -1K$

[5] physical addr:


$28K + (-1K) = 27K$

Address checking:


$|-1K| < 2K$

Visit physical memory:

27K



Segmentation	Base	Size	Grows Postive
Code	32K	2K	1
Heap	34K	2K	1
Stack	28K	2K	0



Segmentation

- Support for Sharing
 - Protection bit

Segmentation	Base	Size	Grows Postive	Protection
Code	32K	2K	1	Read- Execute
Heap	34K	2K	1	Read-Write
Stack	28K	2K	0	Read-Write

Segmentation

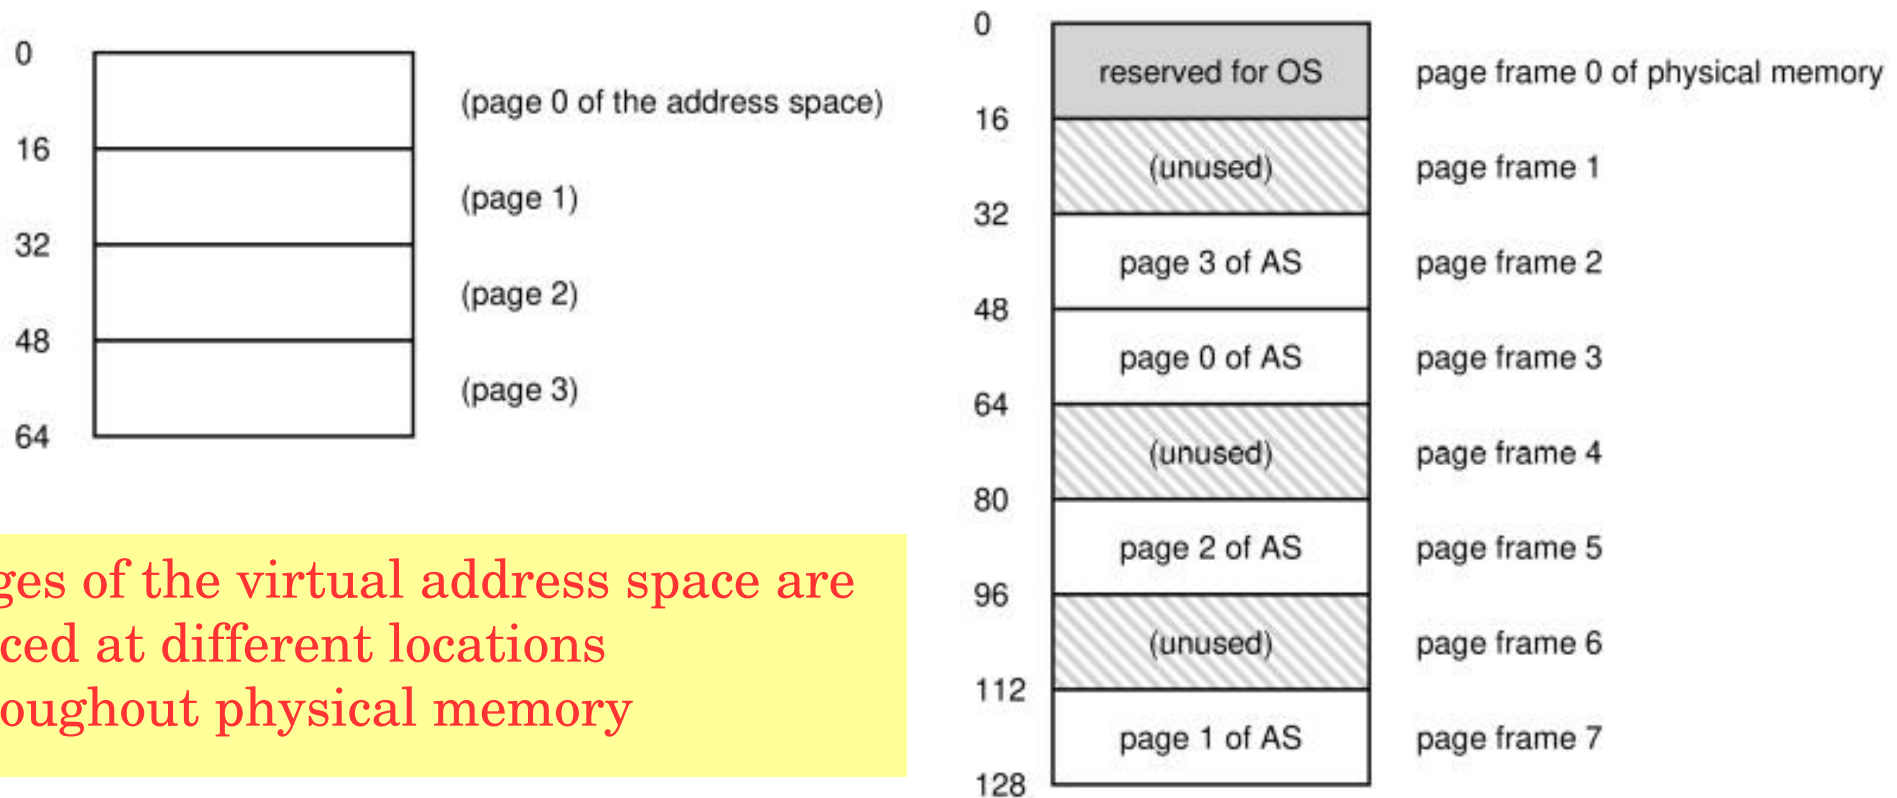
- Summary
 - Base/Bound registers in MMU
 - Multiple Base/Bound
 - Growth direction
 - Protection
- Problem
 - Where to place new address spaces
 - External fragmentation
 - Free memory management

Paging

- Segmentation
 - Splitting address space with variable size logical segmentations
- Paging
 - Divide address space into fixed size units (pages)

Paging

- Example:
 - 64 Byte address space (i.e., 6 bit pointer)
 - 16 Byte page
 - 128 Byte physical memory



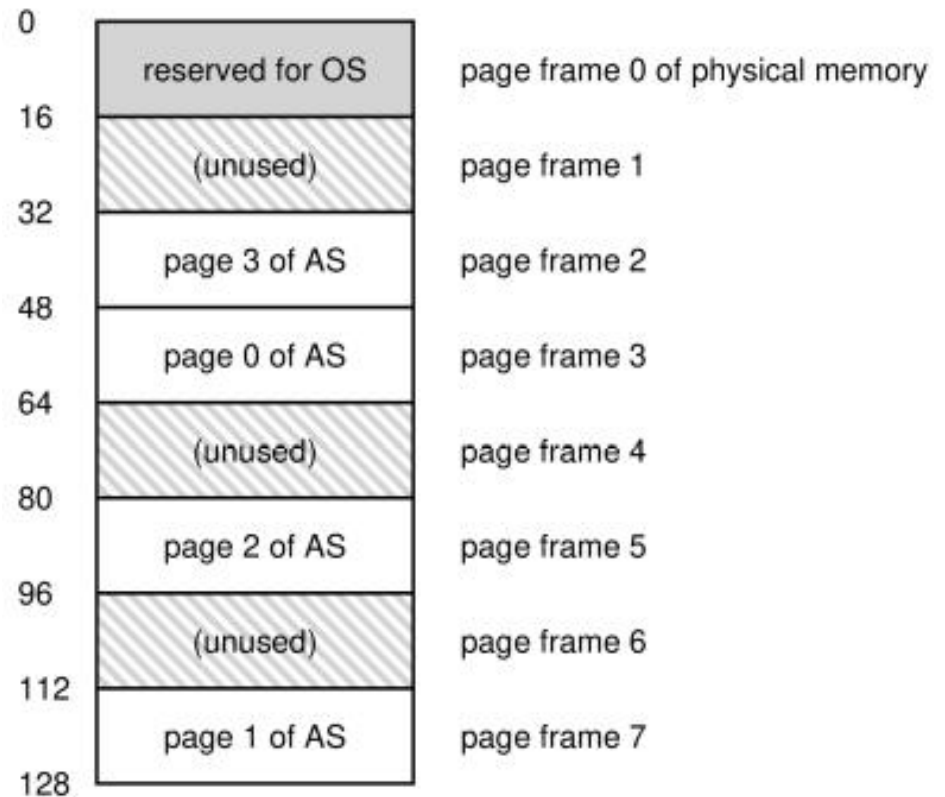
Pages of the virtual address space are placed at different locations throughout physical memory

Paging

- Advantages
 - Flexible
 - make no assumptions about the direction the heap/stack grow, how they are used.
 - Simple
 - Simple free memory management
 - A free list of free pages

Paging

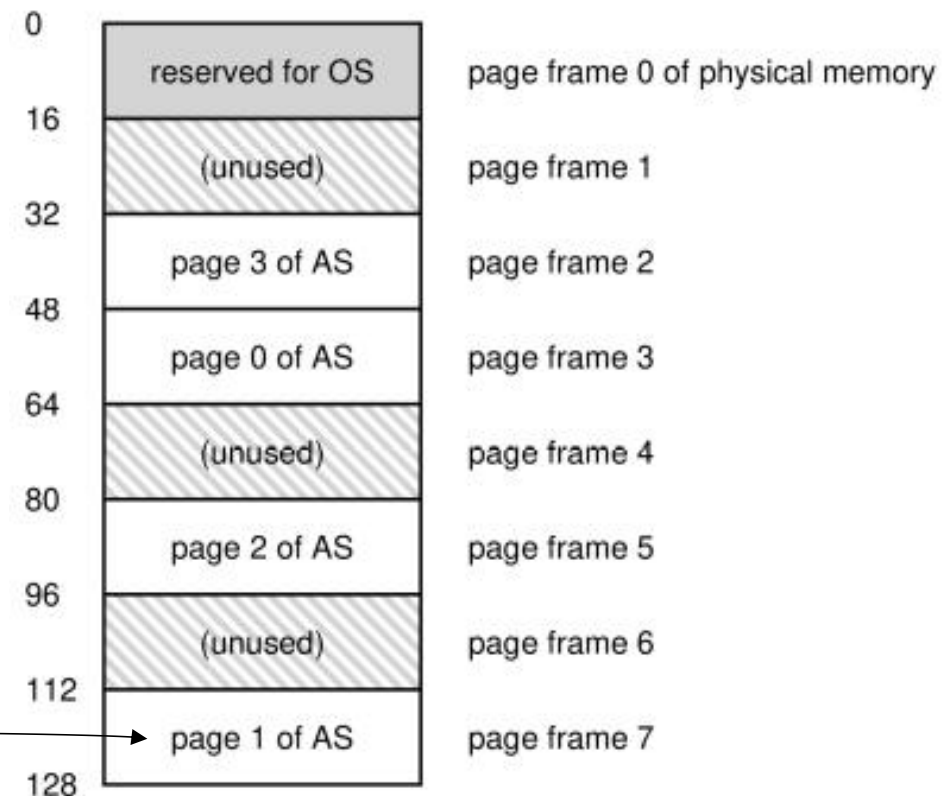
- Virtual page → physical frame
 - **Page Table**
 - A data structure
 - VP0 → PF3
 - VP1 → PF7
 - VP2 → PF5
 - VP3 → PF2
 - In each process



Paging

- Address translation
 - Virtual address:
 - **Virtual Page Num (VPN)**
 - **Offset**
 - Example
 - 64 Byte virtual address (6 bit pointer)
 - 16 Byte per page

- Address translation
 - `movl 21, %eax`
 - Binary of 21: 010101
 - the 5-th byte (0101) of the first virtual page (01)
- VP1 → FP7




Paging

- Questions
 - Where are page tables stored?
 - What are the typical contents of the page table?
 - How big are the tables?
 - Does paging make the system (too) slow?

Paging

- Address translation



Paging

- How big are the tables?
 - 32bit address space
 - 4K page size
 - 20bit VPN + 12bit offset
 - $2^{20} = 1\text{M}$
 - translations that the OS would manage
 - For each process!
- Page Table Entry (PTE)
 - 4 Byte
- Page table size: $2^{20} * 4 = 4\text{M}$
- If we have 100 active processes: 400M
- How about 64bit systems?

Paging

- Where are page tables stored?
 - Not in MMU (so big)
 - In OS's memory
 - Physical memory managed by kernel
 - Virtual memory of OS (can be swapped out)

Paging

- What's actually in a page table?
 - Page Table Entry (PTE)
 - An array (linear page table)
 - OS indexes the array with VPN
- PTE
 - PFN
 - Valid bit: whether the VPN is unused
 - Protection bit: read/write/execute
 - Present bit: whether the page on physical memory or on disk (swapped out)
 - Dirty bit: whether the page has been modified since it is brought into memory
 - Reference bit: whether a page has been accessed

Paging

- Too slow

```
VPN = (VirtualAddress & VPN_MASK) >> SHIFT
```

```
PTEAddr = PageTableBaseRegister + (VPN * sizeof(PTE))
```


- Example

```
int array[1000];  
...  
for (i = 0; i < 1000; i++)  
    array[i] = 0;
```

```
0x1024 movl $0x0, (%edi,%eax,4)  
0x1028 incl %eax  
0x102c cmpl $0x03e8, %eax  
0x1030 jne 0x1024
```

Paging

- Too slow



Paging

- Faster translation
 - With the help of hardware (in MMU)
 - Translation Lookaside Buffer (TLB)
 - Cache
 - Temporal and spatial locality
- Smaller page table
 - Hybrid segmentation and paging
 - Multi-layer page table